

HDS Ledger

Stage 2

Highly Dependable Systems

Alameda

Group 27

Catarina Beirolas
93034

João Luís
95607

14th April 2023

1 Design

The goal of the project's second stage is to create a simplified Smart Contract application that establishes a Token Exchange System (TES). This application utilizes the functionality already developed in the first stage of the project to build a State Machine Replication system that can effectively support the TES, which holds a set of accounts, where each account has the following properties:

- It is uniquely identified by a public key from an asymmetric key pair;
- It holds the current balance.

Because it would be expensive to perform consensus for every single transaction individually, to optimize performance, we opted to define a fixed block size that accommodates a specific number of transactions, 4 in our implementation. Only transactions that modify the state of the blockchain are executed in block, operations that only query the system run a simplified protocol we later describe in more detail.

The system membership is static for the entire system lifetime, including a predefined leader and the membership information is known by all participating processes before the start of the system. We assume every node in the system (servers and clients) might be Byzantine.

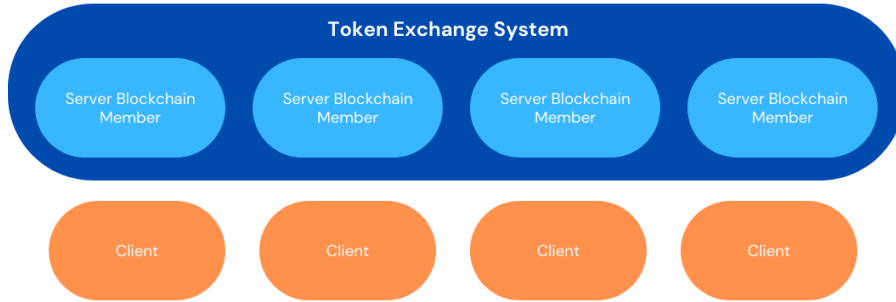


Figure 1: Design Overview

1.1 Client

The Client has a CLI through which it can submit requests to the system. The requests available are: create an account, transfer an amount to other clients and check the account's current balance. These operations are explained in detail in Section 3, Available Operations. The Client always has to first create an account before doing any other operation. For every request, a reply is received confirming if and when the request was executed.

1.2 Server

The Servers are responsible for maintaining the set of TES accounts and ensure a set of dependability properties and security guarantees, namely:

- No account should have a non-negative balance;
- The state of the accounts cannot be modified by unauthorized users;
- Guaranteeing the non-repudiation of all operations issued on an account.

1.2.1 Channels

As required the basic communication is done using **UDP** which allows the network to approximate the behavior of Fair Loss Links. To simulate Perfect Channels, we implement the following:

- To avoid message loss, we set a 5 second timeout and re-send the message until a reply is received;
- To avoid message duplication, we added a messageId to every message;
- To guarantee Integrity and Authentication, our messages are signed with the sender's private key and the receiver verifies it with the sender's public key;
- To guarantee freshness in the communication, the message's sender adds a nonce and then verifies if the nonce in the reply is the expected one.

1.2.2 Broadcast

We chose to use Best Effort Broadcast in our project. To simulate Perfect Channels and guarantees such as freshness, integrity, non-loss and non-duplication we implemented mechanisms which we explain later in the report.

1.2.3 Command Processing

We use threads to allow the Servers to still receive requests from the Clients while a Consensus instance is happening. When the consensus protocol has started, if a Server receives a blockchain modifying message from the Client, it is added to a queue. When Servers are deciding the requests in the block, they remove the first string in the queue, so that all requests from the Client are only processed once and in order.

1.2.4 Command Translation

The Servers translate Client requests to invocations of the Istanbul BFT protocol [1], and also for the respective response to the Client. When a Client makes a request to the service, when the block is filled, the leader server begins the Consensus Protocol with the START procedure of the Algorithm 1. At the end of the Consensus, after the DECIDE step of the Algorithm 2, the Servers answer the Client with an "ACK" message. When the Client has two $(f+1)$ "ACK" messages, it knows the command has been applied.

1.2.5 Leader

The leader doesn't change during the system lifetime and is determined by the Server with the lowest port.

1.3 Keys

We assume there is a Public Key Infrastructure in place. Blockchain members and blockchain clients use public/private keys that we previously generated, which are pre-distributed before the start of the system, and this information is also included in the static system membership.

2 Threats and Protection Mechanisms

In order to guarantee the required security and design specifications, we designed a set of demo tests (puppet masters) that create and run 4 servers and 4 clients.

2.1 Byzantine Server sends multiple ACKs/NACKs to the client for a given request (instead of only once)

A Client waits for a quorum of ACKs/NACKs from the Servers to know if the request was completed or not. A Byzantine Server could try to send multiple ACKs/NACKs to deceive the Client into thinking a quorum was reached. To avoid this threat, the Client checks, for a given requestId if it has already received an ACK/NACK for a given port. If it has, they ignore the message and don't take into account for the quorum.

2.2 Byzantine Server sends repeated "PREPARE" / "COMMIT" messages during the consensus instance trying to forge a quorum

The approach taken is the same as described before: the Servers also check if for a given consensus instance, for a given phase, they have already received a message from a given port. If they have, the request is ignored and not taken into account for the consensus quorum.

2.3 Man in the middle tries to drop messages

All messages exchanged between Clients and Servers (operations or consensus messages) have a timeout: if an answer has not yet been received, it is sent again in an exponential backoff timeout. The ACK/NACK messages are also retransmitted to reduce this threat.

2.4 Byzantine Server send in the block a request signed by them instead of the Client

Since all replicas verify the signatures in all the requests sent in the block, they are able to detect if a Byzantine Server introduced a request there trying to pretend it was a Client request. TODO

2.5 Byzantine Server sends a different value in the StrongReadPhase1

StrongReadPhase1 o byzantino mandar um valor diferente dos outros para testar a Phase2

3 Available Operations

3.1 Create Account

When a Client creates an account, they send their public key in the request, so the Servers verify the Client isn't trying to open an account with a public which is not theirs. If this rule is not met, the Servers answer the Client with a NACK. The Client also has to send a positive amount to open the account. If any of these rules are not met, the Servers answer the Client with a NACK.

3.2 Check Balance

This is an operation that only queries (and does not update) the state of the blockchain, so it is allowed to run a simplified protocol that supports both the following consistency modes.

3.2.1 Strongly consistent read

Allows for reads to execute faster than normal blockchain operations while still obeying the same atomic (linearizable) semantics that are provided by the Istanbul BFT protocol. We implement it by TODO

3.2.2 Weakly consistent read

Allows for reads to provide a correct but stale output, in a way that such reads can still execute in weak connectivity scenarios, namely when the client can only connect to a single replica. We implement it using TODO

3.3 Transfer

A Client can request a transfer of a positive amount between a pair of accounts provided:

- They own the corresponding private key of the source account where the money is withdrawn from (and so have authority to perform the transfer);
- They have enough balance in their account to cover the amount to be transferred plus the fee paid to the leader (block producer);
- The destination account exists;

If any of these requirements are not met, the Servers answer the Client with a NACK.

4 Dependability Guarantees and Security Attributes

Here we'll discuss some of the essential characteristics related to security and dependability and explain why we included them and how we achieved them.

Confidentiality: based on the project requirements, it was decided that confidentiality wouldn't be necessary for this implementation.

Integrity: since we are using UDP, the network is unreliable and communication channels are not secured, every time any data is transferred from Client to Server and between Servers, we perform integrity checks through the usage of message signatures. This means that any time we send a message, we hash it (using SHA1) and encrypt it with the sender's private key. Then the receiver must decrypt the message with the sender's public key and verify if the sender is correct.

Availability: a subset of the blockchain members may be malicious and behave in an arbitrary manner. However, as long as the Client receives $f+1$ "ACK" messages, the service is available.

Reliability: we can provide reliability over the period of time that the servers are active and correct, because we still need to reach a quorum.

Freshness: in the communication between the Client and the Servers and between Servers, the messages are sent with a message nonce (which is also signed with the sender's private key) and when a reply is received, the sender checks if the message nonce received corresponds to the expected value, guaranteeing it is not an old message being replayed.

Authentication: our project provides integrity and freshness, so it also guarantees authentication.

4.1 References

[1] Henrique Moniz. The Istanbul BFT Consensus Algorithm.

<https://arxiv.org/pdf/2002.03613.pdf>